
Lab 2: Wikipedia

MSE - Big Data Analytics (BDA) 2019

Pedro Costa, Thibaut Piquerez, Jonathan Guerne



MASTER OF SCIENCE
IN ENGINEERING

2019-04-04

Contents

1	Introduction	3
2	Implementation	3
2.1	Part 1 RankLangs	4
2.2	Part 2 RankLangsUsingIndex	5
2.3	Part 3 RankLangsReduceByKey	8
2.4	Performance	9
3	Conclusion	10

1 Introduction

The aim of this laboratory is to discover the framework spark. This will be done in a practical case as we will use scala and spark to rank programming language over based on their **popularity**. To be more precise we will try and compare different methods to compute the number of time a specific language was mentioned in a series of Wikipedia articles.

In the conclusion of this report we will also compare our computed results with the ones presented in this recent article from the RedMonk. RedMonk computed their statistics mostly with Github and stackoverflow so it will be interesting to take that in account while comparing the results.

Although spark offers the ability to be used in a distributed system this laboratory will be implemented with one only node (our laptop).

The list of wikipedia articles were given with the laboratory instructions.

2 Implementation

As mentioned before the goal of this laboratory is to compare 3 methods of computing the number of time a specific language was mentioned throughout a list of wikipedia articles. This laboratory will be done in scala using the framework spark. Our IDE is IntelliJ by JetBrains.

A series of Wikipedia articles were given with this laboratory. We will use these articles to compute our results. With the articles the laboratory also comes with helpful code to make the interaction with the articles easier. We will not describe this given code in our report.

In part 1 we will implement a naive approach of the language ranking by dividing our problem in two parts. First computing the number of occurrences of a specific language throughout all articles and then secondly, using our previously built function to rank all languages based on their number of mentions.

To assess that our implementation is working correctly this laboratory also comes with a series of unit tests (multiple for each part). We were also invited, if we wanted too, to add tests we found relevant but we didn't find the need to.

2.1 Part 1 RankLangs

```
1 def occurrencesOfLang(lang: String, rdd: RDD[WikipediaArticle]): Int =  
  {  
2   rdd.aggregate(0)((acc, article) =>  
3     if (wordInArticle(lang, article))  
4       acc + 1  
5     else  
6       acc,  
7     (x, y) => x + y)  
8  }
```

As displayed here the first part of the code doesn't take much space. It is a rather simplistic approach which will increment the accumulator if the condition `WordInArticle` is true. The `WordInArticle` function look as follow :

```
1 def wordInArticle(w: String, a: WikipediaArticle): Boolean = {  
2   a.text.split("[ ,!.]+" ).contains(w)  
3 }
```

This function is really time consuming compared to more simplistic approach (see more about that in the **Performance** chapter). But it is not complicated. It takes a `WikipediaArticle` and a language, split the content of the article into a word array and then check if this word array contains the specified language.

Finally to compute the ranking we call one last function

```
1 def rankLangs(langs: List[String], rdd: RDD[WikipediaArticle]): List[(  
  String, Int)] = {  
2   langs.map((l: String) => (l, occurrencesOfLang(l, rdd))).sortBy(_._2).  
    reverse  
3 }
```

For each language in the list we will compute the number of mention throughout all articles and generate a tuple (`language`, `# of mentions of the language`). We can then sort all theses tuples based on their second elements and reverse the list to obtain a descending sorted list of language and the number of times they were mentioned in articles.

We verified our implementation of this part by launching the given unit test. All of them passed.

How much does the code take? What is the list of ranked languages?

The total computation time for this approach was :

```
1 Processing Part 1: naive ranking took 59018 ms.
```

So around 1 minute witch is really long and should be improved.

It output the following list :

Language	Number of mention
JavaScript	1704
C#	731
Java	699
CSS	429
Python	409
C++	384
PHP	333
MATLAB	296
Perl	175
Ruby	160
Haskell	65
Objective-C	61
Scala	53
Clojure	29
Groovy	29

It it good to see that this table is similar in each implemented approach

2.2 Part 2 RankLangsUsingIndex

This approach is divided in two part. We will first implement an inverted index solution witch will be able to return a `RDD` of tuples `(String, Iterable[WikipediaArticle])`. The idea of inverted index is that for each element (in our case for each language) we will have a list of article mentioning it. It is really interesting because it allow fast full text searches.

```
1  def makeIndex(langs: List[String], rdd: RDD[WikipediaArticle]): RDD[(  
    String, Iterable[WikipediaArticle])] = {  
2  
3      val arts = rdd.flatMap((a: WikipediaArticle) =>  
4          a.text.split("[ ,!.]+" )  
5              .toSet  
6              .filter((word: String) =>  
7                  langs.contains(word))  
8              .map((_, a)))  
9          .collect()  
10  
11      sc.parallelize(arts).groupByKey()  
12  
13  }
```

Once this inverted index build we can use it to compute the ranking.

```
1  def rankLangsUsingIndex(index: RDD[(String, Iterable[WikipediaArticle  
    ])]): List[(String, Int)] = {  
2      index.mapValues((it: Iterable[WikipediaArticle]) => it.size).  
          sortBy(_._2).collect().toList  
3  }
```

By using `MapValue` we are able to compute quickly the number of articles mentioning a certain language. We then only have to sort the computed list.

How much does the code take? What is the list of ranked languages?

The total computation time for this approach was :

```
1  Processing Part 2: ranking using inverted index took 11732 ms.
```

Witch is already way faster that in part 1.

It output the following list :

Language	Number of mention
JavaScript	1704
C#	731
Java	699
CSS	429
Python	409

Language	Number of mention
C++	384
PHP	333
MATLAB	296
Perl	175
Ruby	160
Haskell	65
Objective-C	61
Scala	53
Clojure	29
Groovy	29

It is good to see that this table is similar in each implemented approach

Can you notice a performance improvement over attempt #1? Why?

There is indeed a significant performance improvement. One of the reasons might be the following : for the first part to check if there was a mention of a specific language in an article we iterated through each language and then used the function `occurrenceOfLang` over each article for a specific language. And it is important to note that to be effective the “check function” has to iterate through each word of a given article’s text.

Knowing all of that it is then pretty straight forward to realize that if we iterate first through the languages and then only through the words in the article will be less cost efficient in terms of list generation than the opposite.

2.3 Part 3 RankLangsReduceByKey

```
1  def rankLangsReduceByKey(langs: List[String], rdd: RDD[
    WikipediaArticle]): List[(String, Int)] = {
2
3  rdd.flatMap((a: WikipediaArticle) =>
4    a.text.split("[ ,!.]+" )
5    .toSet
6    .filter((word: String) =>
7      langs.contains(word))
8    .map((_, 1)))
9    .reduceByKey(_ + _)
10   .sortBy(-_._2)
11   .collect()
12   .toList
13
14 }
```

How much does the code take? What is the list of ranked languages?

The total computation time for this approach was :

```
1 Processing Part 3: ranking using reduceByKey took 6381 ms.
```

Witch is by far the fastest approach.

It output the following list :

Language	Number of mention
JavaScript	1704
C#	731
Java	699
CSS	429
Python	409
C++	384
PHP	333
MATLAB	296
Perl	175
Ruby	160

Language	Number of mention
Haskell	65
Objective-C	61
Scala	53
Clojure	29
Groovy	29

It is good to see that this table is similar in each implemented approach

Can you notice an improvement in performance compared to measuring both the computation of the index and the computation of the ranking as we did in attempt #2? If so, can you think of a reason?

2.4 Performance

- 1 Processing Part 1: naive ranking took 59018 ms.
- 2 Processing Part 2: ranking using inverted index took 11732 ms.
- 3 Processing Part 3: ranking using reduceByKey took 6381 ms.

TODO : mentionner le problème du temps d'exécution avec le contain et les améliorations

The plot below will summarize our results :




Figure 1: Languages mentioned in Wikipedia articles

3 Conclusion

Rank	RedMonk analysis	Our Analysis
1	Javascript	Javascript
2	Java	C#
3	Python	Java
4	PHP	CSS
5	C#	Python
6	C++	C++
7	CSS	PHP
8	Ruby	MATLAB
9	C	Perl
10	Objective-C	Ruby
11	Swift	Haskell
12	Scala	Objective-C
13	Shell	Scala

Rank	RedMonk analysis	Our Analysis
14	Go	Clojure
15	R	Groovy
16	TypeScript	-
17	Powershell	-
18	Perl	-
19	Haskell	-
20	Lua	-

For the most part we can see that there is a slight different between the two approach although we can find some exact similarities (Javascript, C++) and some pretty close call (Java, Python, Objective-C).

Firstly, To understand this difference we may have to understand the difference between the medium used to compute our results. As said before the RedMonk analysis was mostly based on Github and Stackoverflow and, obviously, our studies was based on Wikipedia's articles. It is clear that a developer will find most of the information he is searching on Stackoverflow/Github rather than Wikipedia. It's also more probable to find older informations on Wikipedia as it's main purpose is to work as an encyclopedia.

Secondly, we can obviously see that our approach used a couple less languages then RedMonk did. It would have been interesting to see how the bottom of the table will have look like if we could have had more information. Our guess is that as languages get less and less mentioned difference between the two approach would have increase.

So, we were able to compute an estimation of the usage of common programming languages. We were also able to compare our work with the analysis made by RedMonk and were able to conclude that although our results differ it may have come from the source of the data itself. We can conclude that by a great margin [JavaScript](#) is today the most popular programming language.